


<u>Concepts</u>	<u>Technologies</u>	<u>Application</u>
Stateless & Stateful	Kubernetes	
Virtualization v/s Containers		
Orchestration		
DevOps		
CI/CD		

Docker

Dockerfile	Multi-host
Docker registry	Cross-host networking
	State/Volume

1. Image building
2. Container ↔ Container networking
3. Volume management






Anytime, anywhere, any device

State, Stateless, Stateful



	Scalability	Failure Handling
Stateless	LB + Replicas	Replace replica
Stateful	Sharding	Replication

Stateless Scalability




Load balancer



1. Ingress
2. Load balancing / Round Robin / Reverse Proxy
3. Failure management

Health check endpoint

/ping → 200 OK




Stateful scalability



Read
CRUD
Write


Types of failure

1. M/c failure $\begin{matrix} \text{S/W} \\ \text{H/W} \end{matrix}$ → Multiple m/c
2. Power, Fire → Multiple A/z
3. Cyclone, Earthquake → Multiple region




Shard key → Consistent Hashing → 0/1


Virtualization



Real Machine



Isolation of resources



AWS → 2006

IaaS

Public

Google Cloud → 2008

PaaS

Private

ELB, RDS → 2009

IAC

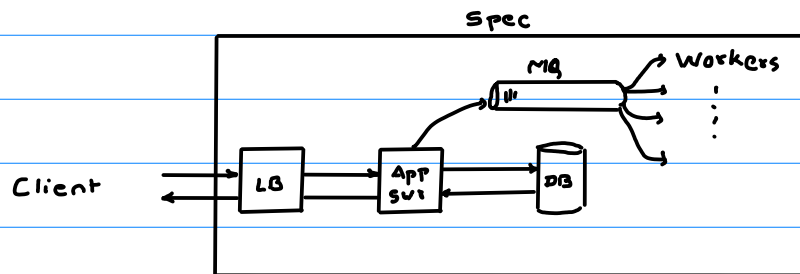
Reduced admin
Increased automation

Hybrid

CloudFormation → 2010

Multi-cloud

Azure → 2010



App Servers

Databases

Load balancers

Caching Servers

Message queues



Vendor Lock-in

Kubernetes → FOSS

↳ Container based orchestration
↳ Google Kubernetes Engine

Amazon Kubernetes Service


Azure — " — " —

Docker

RedHat / IBM

VMware / Pivotal

Containers

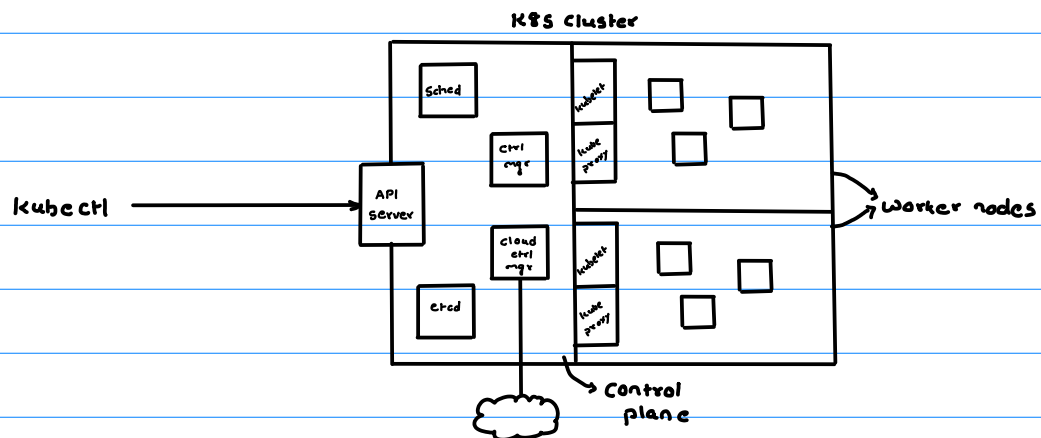



cgroups
namespaces
chroot

Containers v/s VM

1. Resource isolation
2. VMs need more resources
Containers are light-weight
3. VMs are more secure
Containers are not as secure

Kubernetes





kubectl get nodes

KUBERNETES / Context

1. Cluster (IP:Port)
2. Auth info
3. Namespace

Resources

Pods	Persistent Volumes	Configmaps
Replicaset	Persistent Volume Claims	Secrets
Deployment	Statefulsets	Namespaces
Services		

Declarative v/s Imperative Spec

deploy.sh

```

#!/bin/bash
kubectl run ---
kubectl run ---
---
```

\$./deploy.sh

- * What to do
- * When to do
- * When not to do

Declarative


What we need → Orchestrator → Env
(Desired state) (Current state)

```

Pods:
1. nginx → image/nginx
2. mysql → ---
---
```

YAML

YAML



Array of arrays → `[[], []]`



Array of objects → `[{ }, { }]`

Object containing arrays → `{ \" \" : [] }`


Object containing objects → `{ \" \" : { } }`

Pods & containers

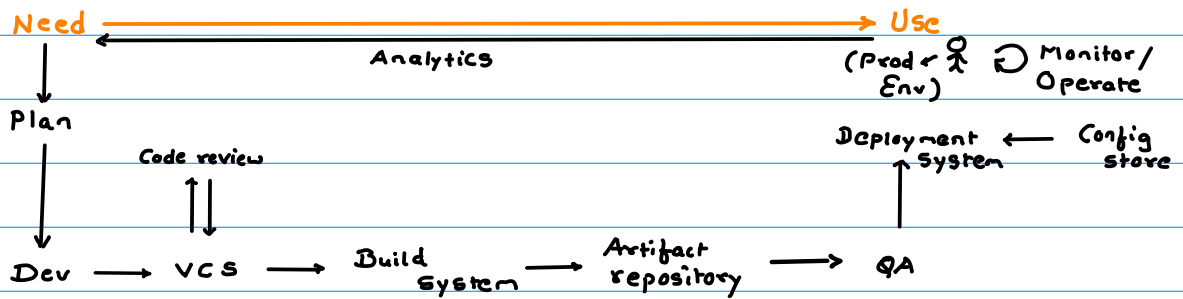
1. All containers of a pod have the same IP
2. All containers of a pod have to be healthy
3. All containers of a pod will be in the same node



Sidecar containers



DevOps




Provisioning

Deployment

v1 v1 v1 → v2 v2 v2

1. Bring down v1
Start all v2 (Recreate)
2. Start all v2
Switch traffic
Bring down v1
3. Incremental deployment
(Rolling update)



Current Desired
v1 v1 v1 → T₀ → v2 v2 v2


✗ v1 v1 v2 → T₁

✗ v1 v2 v2 → T₂






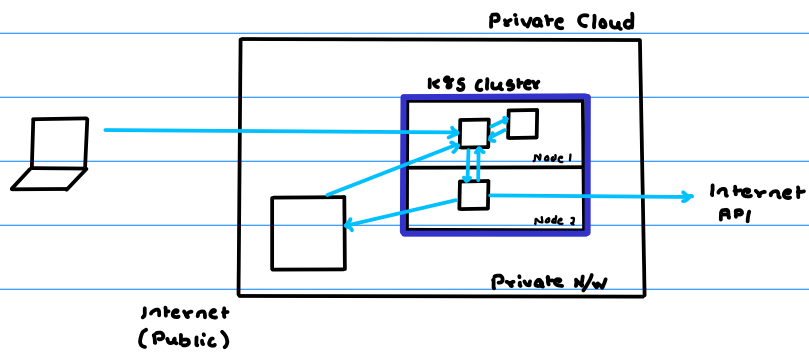
✗ v2 v2 v2 → T₃

v2 v2 v2 → T₄


4. Canary deployment



Services



1. Service
2. Kube DNS nginx → 10.108.215.91
3. Network driver Node ↔ Node
4. kube-proxy
5. endpoint-controller



/users → users:80



/payments → payments:80

↳ Cluster IP : Port

Volume/Storage Management

MySQL → DDL & DML

Data directory
(/var/lib/mysql)




docker run -v /data/disk1 : /var/lib/mysql



Host path

Container path

Write → Read

Read ← Write





1. Create Volume

2. Attach volume → /dev/sd**f**

3. Partition volume → **f**disk → /dev/sd**f**1
/dev/sd**f**2

4. Format partition → mkfs.*

5. Mount partition → mount /dev/sd**f**1 /mnt/disk1

6. Read/write file/directories → cat, ls, rm, mkdir ...

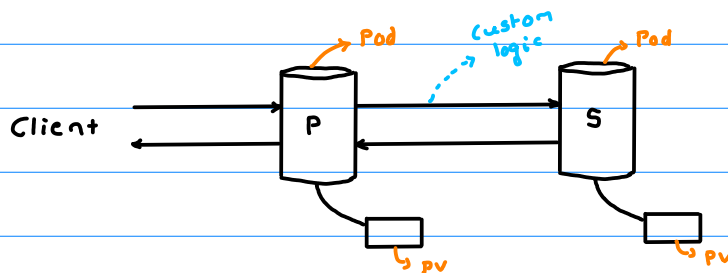

7. Umount → umount

8. Detach volume

Cloud APIs


O/S

Persistent Volume (PV) & Persistent Volume Claim (PVC)



<u>Primary</u>	<u>Secondary</u>	<u>Secondary</u>
Pod	Pod	Pod
PVC	PVC	PVC
PV	PV	PV
Service	Service	Service

Service discovery in stateful components



Client needs to know the specific pod
Pods need to communicate with each other

Replicaset

replicas: 3
template:
pod spec




Pod names were
not predictable

StatefulSet

replicas: 3
template:
pod spec
vc-template:
pvc spec

4. Every pod can be addressed



Cluster Role (non-namespaced resources)

